

# Customer Segmentation with K - Means Clustering

## Project Objective

The goal of this project is to perform customer segmentation using a dataset of customer demographics, spending patterns, and purchase behavior.

By applying clustering techniques, we aim to identify distinct behavioral personas that can help businesses better understand their customers and design targeted marketing strategies.

## Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv("marketing_campaign.csv", sep="\t")
df
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Comp
0	5524	1957	Graduation	Single	58138.0	0	0	04
1	2174	1954	Graduation	Single	46344.0	1	1	08
2	4141	1965	Graduation	Together	71613.0	0	0	21
3	6182	1984	Graduation	Together	26646.0	1	0	10
4	5324	1981	PhD	Married	58293.0	1	0	19
...	...	...	...	...	...	...	...	...
2235	10870	1967	Graduation	Married	61223.0	0	1	13
2236	4001	1946	PhD	Together	64014.0	2	1	10
2237	7270	1981	Graduation	Divorced	56981.0	0	0	25
2238	8235	1956	Master	Together	69245.0	0	1	24
2239	9405	1954	PhD	Married	52869.0	1	1	15

2240 rows × 29 columns

## Data Preprocessing

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ID               2240 non-null    int64  
 1   Year_Birth       2240 non-null    int64  
 2   Education        2240 non-null    object  
 3   Marital_Status   2240 non-null    object  
 4   Income            2216 non-null    float64 
 5   Kidhome          2240 non-null    int64  
 6   Teenhome         2240 non-null    int64  
 7   Dt_Customer      2240 non-null    object  
 8   Recency           2240 non-null    int64  
 9   MntWines          2240 non-null    int64  
 10  MntFruits         2240 non-null    int64  
 11  MntMeatProducts  2240 non-null    int64  
 12  MntFishProducts  2240 non-null    int64  
 13  MntSweetProducts 2240 non-null    int64  
 14  MntGoldProds     2240 non-null    int64  
 15  NumDealsPurchases 2240 non-null    int64  
 16  NumWebPurchases  2240 non-null    int64  
 17  NumCatalogPurchases 2240 non-null    int64  
 18  NumStorePurchases 2240 non-null    int64  
 19  NumWebVisitsMonth 2240 non-null    int64  
 20  AcceptedCmp3     2240 non-null    int64  
 21  AcceptedCmp4     2240 non-null    int64  
 22  AcceptedCmp5     2240 non-null    int64  
 23  AcceptedCmp1     2240 non-null    int64  
 24  AcceptedCmp2     2240 non-null    int64  
 25  Complain          2240 non-null    int64  
 26  Z_CostContact    2240 non-null    int64  
 27  Z_Revenue          2240 non-null    int64  
 28  Response           2240 non-null    int64  
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB
```

In [4]: `df.isnull().sum()`

```
Out[4]: ID          0  
Year_Birth      0  
Education       0  
Marital_Status  0  
Income          24  
Kidhome         0  
Teenhome        0  
Dt_Customer     0  
Recency         0  
MntWines        0  
MntFruits       0  
MntMeatProducts 0  
MntFishProducts 0  
MntSweetProducts 0  
MntGoldProds    0  
NumDealsPurchases 0  
NumWebPurchases 0  
NumCatalogPurchases 0  
NumStorePurchases 0  
NumWebVisitsMonth 0  
AcceptedCmp3    0  
AcceptedCmp4    0  
AcceptedCmp5    0  
AcceptedCmp1    0  
AcceptedCmp2    0  
Complain        0  
Z_CostContact   0  
Z_Revenue        0  
Response         0  
dtype: int64
```

```
In [5]: df.duplicated().sum()
```

```
Out[5]: np.int64(0)
```

The missing values in the income column represent about 1% of the dataset, so dropping them will not significantly reduce the data size.

```
In [6]: df = df.dropna()
```

```
In [7]: df.nunique()
```

```
Out[7]:   ID          2216
          Year_Birth    59
          Education      5
          Marital_Status  8
          Income         1974
          Kidhome        3
          Teenhome       3
          Dt_Customer    662
          Recency        100
          MntWines       776
          MntFruits      158
          MntMeatProducts 554
          MntFishProducts 182
          MntSweetProducts 176
          MntGoldProds    212
          NumDealsPurchases 15
          NumWebPurchases 15
          NumCatalogPurchases 14
          NumStorePurchases 14
          NumWebVisitsMonth 16
          AcceptedCmp3     2
          AcceptedCmp4     2
          AcceptedCmp5     2
          AcceptedCmp1     2
          AcceptedCmp2     2
          Complain         2
          Z_CostContact    1
          Z_Revenue         1
          Response         2
          dtype: int64
```

The columns `Z_CostContact` and `Z_Revenue` have only 1 unique value, meaning they do not provide any useful information for analysis or modeling. Therefore, we can safely remove them from the dataset.

```
In [8]: df = df.drop(columns=['Z_CostContact', 'Z_Revenue'])
```

The `Dt_Customer` column was converted to datetime using `.loc` to avoid pandas warnings, ensuring proper handling of date operations.

```
In [9]: df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'].astype(str), errors='coerce')
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2216 entries, 0 to 2239
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ID               2216 non-null    int64  
 1   Year_Birth       2216 non-null    int64  
 2   Education        2216 non-null    object  
 3   Marital_Status   2216 non-null    object  
 4   Income            2216 non-null    float64 
 5   Kidhome          2216 non-null    int64  
 6   Teenhome         2216 non-null    int64  
 7   Dt_Customer      2216 non-null    datetime64[ns]
 8   Recency           2216 non-null    int64  
 9   MntWines          2216 non-null    int64  
 10  MntFruits         2216 non-null    int64  
 11  MntMeatProducts  2216 non-null    int64  
 12  MntFishProducts  2216 non-null    int64  
 13  MntSweetProducts 2216 non-null    int64  
 14  MntGoldProds     2216 non-null    int64  
 15  NumDealsPurchases 2216 non-null    int64  
 16  NumWebPurchases  2216 non-null    int64  
 17  NumCatalogPurchases 2216 non-null    int64  
 18  NumStorePurchases 2216 non-null    int64  
 19  NumWebVisitsMonth 2216 non-null    int64  
 20  AcceptedCmp3     2216 non-null    int64  
 21  AcceptedCmp4     2216 non-null    int64  
 22  AcceptedCmp5     2216 non-null    int64  
 23  AcceptedCmp1     2216 non-null    int64  
 24  AcceptedCmp2     2216 non-null    int64  
 25  Complain          2216 non-null    int64  
 26  Response          2216 non-null    int64  
dtypes: datetime64[ns](1), float64(1), int64(23), object(2)
memory usage: 484.8+ KB
```

```
In [11]: newest_date = df['Dt_Customer'].max()
oldest_date = df['Dt_Customer'].min()

print(f"The newest customer's enrolment date in the records: {newest_date.date()}")
print(f"The oldest customer's enrolment date in the records: {oldest_date.date()})
```

The newest customer's enrolment date in the records: 2014-06-29  
The oldest customer's enrolment date in the records: 2012-07-30

The Customer\_Tenure column was created by calculating the number of days between each customer's enrolment date and the newest enrolment date in the records.

```
In [12]: reference_date = df['Dt_Customer'].max()
df.loc[:, 'Customer_Tenure'] = (reference_date - df['Dt_Customer']).dt.days
```

```
In [13]: df['Marital_Status'].value_counts()
```

```
Out[13]: Marital_Status
Married      857
Together     573
Single       471
Divorced     232
Widow        76
Alone         3
Absurd        2
YOLO          2
Name: count, dtype: int64
```

The 'Alone' category in Marital\_Status will be replaced with 'Single' to consolidate similar categories and reduce noise in the clustering process.

```
In [14]: df.loc[:, 'Marital_Status'] = df['Marital_Status'].replace({'Alone': 'Single'})
```

Rows with rare and inconsistent categories 'Absurd' and 'YOLO' in Marital\_Status will be removed, as they represent a negligible portion of the dataset

```
In [15]: df=df[~df['Marital_Status'].isin(['Absurd', 'YOLO'])]
```

```
In [16]: df['Marital_Status'].value_counts()
```

```
Out[16]: Marital_Status
Married      857
Together     573
Single       474
Divorced     232
Widow        76
Name: count, dtype: int64
```

```
In [17]: df['Education'].value_counts()
```

```
Out[17]: Education
Graduation    1115
PhD           479
Master         364
2n Cycle      200
Basic          54
Name: count, dtype: int64
```

The Age column will be created by subtracting Year\_Birth from 2025, providing the age of each customer for further analysis and clustering.

```
In [18]: df.loc[:, 'Age'] = 2025 - df['Year_Birth']
```

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2212 entries, 0 to 2239
Data columns (total 29 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ID               2212 non-null    int64  
 1   Year_Birth       2212 non-null    int64  
 2   Education        2212 non-null    object  
 3   Marital_Status   2212 non-null    object  
 4   Income            2212 non-null    float64 
 5   Kidhome          2212 non-null    int64  
 6   Teenhome         2212 non-null    int64  
 7   Dt_Customer      2212 non-null    datetime64[ns]
 8   Recency           2212 non-null    int64  
 9   MntWines          2212 non-null    int64  
 10  MntFruits         2212 non-null    int64  
 11  MntMeatProducts  2212 non-null    int64  
 12  MntFishProducts  2212 non-null    int64  
 13  MntSweetProducts 2212 non-null    int64  
 14  MntGoldProds     2212 non-null    int64  
 15  NumDealsPurchases 2212 non-null    int64  
 16  NumWebPurchases  2212 non-null    int64  
 17  NumCatalogPurchases 2212 non-null    int64  
 18  NumStorePurchases 2212 non-null    int64  
 19  NumWebVisitsMonth 2212 non-null    int64  
 20  AcceptedCmp3     2212 non-null    int64  
 21  AcceptedCmp4     2212 non-null    int64  
 22  AcceptedCmp5     2212 non-null    int64  
 23  AcceptedCmp1     2212 non-null    int64  
 24  AcceptedCmp2     2212 non-null    int64  
 25  Complain          2212 non-null    int64  
 26  Response           2212 non-null    int64  
 27  Customer_Tenure   2212 non-null    int64  
 28  Age               2212 non-null    int64  
dtypes: datetime64[ns](1), float64(1), int64(25), object(2)
memory usage: 518.4+ KB
```

In [20]: `df['Kidhome'].value_counts()`

Out[20]: Kidhome

0	1279
1	887
2	46

Name: count, dtype: int64

In [21]: `df['Teenhome'].value_counts()`

Out[21]: Teenhome

0	1145
1	1016
2	51

Name: count, dtype: int64

The Children column will be created by summing Kidhome and Teenhome, providing the total number of children in each customer's household.

In [22]: `df.loc[:, 'Children'] = df['Kidhome'] + df['Teenhome']`

```
In [23]: df['Children'].value_counts()
```

```
Out[23]: Children
1    1115
0     631
2     416
3      50
Name: count, dtype: int64
```

The Children column shows that most customers have 0 or 1 child/teen, while a smaller portion has 2 or 3, representing the total children per household.

```
In [24]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2212 entries, 0 to 2239
Data columns (total 30 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ID               2212 non-null    int64  
 1   Year_Birth       2212 non-null    int64  
 2   Education        2212 non-null    object  
 3   Marital_Status   2212 non-null    object  
 4   Income            2212 non-null    float64 
 5   Kidhome          2212 non-null    int64  
 6   Teenhome         2212 non-null    int64  
 7   Dt_Customer      2212 non-null    datetime64[ns]
 8   Recency          2212 non-null    int64  
 9   MntWines          2212 non-null    int64  
 10  MntFruits         2212 non-null    int64  
 11  MntMeatProducts  2212 non-null    int64  
 12  MntFishProducts  2212 non-null    int64  
 13  MntSweetProducts 2212 non-null    int64  
 14  MntGoldProds     2212 non-null    int64  
 15  NumDealsPurchases 2212 non-null    int64  
 16  NumWebPurchases  2212 non-null    int64  
 17  NumCatalogPurchases 2212 non-null    int64  
 18  NumStorePurchases 2212 non-null    int64  
 19  NumWebVisitsMonth 2212 non-null    int64  
 20  AcceptedCmp3     2212 non-null    int64  
 21  AcceptedCmp4     2212 non-null    int64  
 22  AcceptedCmp5     2212 non-null    int64  
 23  AcceptedCmp1     2212 non-null    int64  
 24  AcceptedCmp2     2212 non-null    int64  
 25  Complain          2212 non-null    int64  
 26  Response           2212 non-null    int64  
 27  Customer_Tenure   2212 non-null    int64  
 28  Age                2212 non-null    int64  
 29  Children           2212 non-null    int64  
dtypes: datetime64[ns](1), float64(1), int64(26), object(2)
memory usage: 535.7+ KB
```

```
In [25]: df.describe()
```

Out[25]:

	ID	Year_Birth	Income	Kidhome	Teenhome	Dt_Cu
<b>count</b>	2212.000000	2212.000000	2212.000000	2212.000000	2212.000000	
<b>mean</b>	5587.731917	1968.811031	52232.510850	0.442586	0.505425	2013-12-17 23:28:06.075
<b>min</b>	0.000000	1893.000000	1730.000000	0.000000	0.000000	2012-01-01 00:00:00
<b>25%</b>	2814.750000	1959.000000	35233.500000	0.000000	0.000000	2013-01-01 00:00:00
<b>50%</b>	5458.500000	1970.000000	51381.500000	0.000000	0.000000	2013-01-01 00:00:00
<b>75%</b>	8421.750000	1977.000000	68522.000000	1.000000	1.000000	2013-01-01 00:00:00
<b>max</b>	11191.000000	1996.000000	666666.000000	2.000000	2.000000	2014-01-01 00:00:00
<b>std</b>	3247.944128	11.982065	25187.455359	0.537052	0.544258	

8 rows × 28 columns



In [26]: `df.nunique()`


Out[26]:

ID	2212
Year_Birth	59
Education	5
Marital_Status	5
Income	1973
Kidhome	3
Teenhome	3
Dt_Customer	662
Recency	100
MntWines	776
MntFruits	158
MntMeatProducts	554
MntFishProducts	182
MntSweetProducts	176
MntGoldProds	211
NumDealsPurchases	15
NumWebPurchases	15
NumCatalogPurchases	14
NumStorePurchases	14
NumWebVisitsMonth	16
AcceptedCmp3	2
AcceptedCmp4	2
AcceptedCmp5	2
AcceptedCmp1	2
AcceptedCmp2	2
Complain	2
Response	2
Customer_Tenure	662
Age	59
Children	4
dtype:	int64


```
In [27]: df['TotalSpent'] = df[['MntWines','MntFruits','MntMeatProducts',
                           'MntFishProducts','MntSweetProducts','MntGoldProds']].sum
```

## Exploratory Data Analysis (EDA)


```
In [28]: plt.figure(figsize=(7,5))
sns.histplot(df['Age'], bins=30, kde=True, color='skyblue')
plt.title('Age distribution')
plt.show()
```



```
In [29]: plt.figure(figsize=(6,4))
sns.countplot(x='Marital_Status', data=df, palette='pastel')
plt.title('Marital_Status Distribution')
plt.show()
```




```
In [30]: plt.figure(figsize=(6,4))
sns.countplot(x='Education', data=df, palette='muted')
plt.title('Education Distribution')
plt.xticks(rotation=45)
plt.show()
```



```
In [31]: plt.figure(figsize=(7,5))
sns.histplot(df['Income'], bins=30, kde=True, color='orange')
```


```
plt.title('Income Distribution')
plt.show()
```



The income distribution is right-skewed, with most customers earning between 20,000 and 80,000. There are a few customers with very high incomes (over \$120,000), but these are outliers. This skew suggests that while the majority of customers have moderate income, targeting high-income individuals could focus on a small, high-value segment.

```
In [32]: plt.figure(figsize=(6,4))
sns.scatterplot(data=df, x='Age', y='Income', hue='Marital_Status', alpha=0.7, p
plt.title('Income vs Age by Marital Status')
plt.show()
```

### Income vs Age by Marital Status



```
In [33]: total_spent=[ 'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts', 'MntSweet'


n_cols = 2
n_rows = -(~len(total_spent)) // n_cols

fig, axes = plt.subplots(n_rows, n_cols, figsize=(6 * n_cols, 6 * n_rows))
axes = axes.flatten()

for i, feature in enumerate(total_spent):
    sns.histplot(data=df, x=feature, hue='Marital_Status', kde=True, bins=30, ax=axes[i])
    axes[i].set_title(f"Distribution of {feature}")

plt.tight_layout()
plt.suptitle("Total spent Based on Marital Status", fontsize=22, y=1.02)
plt.show()
```

### Total spent Based on Marital Status



The histograms show the distribution of spending on different product categories segmented by Marital Status. Customers who are married or together tend to spend more on most categories, especially on wines and meat products. Single, divorced, or

widowed customers generally spend less, with fewer high spenders. This indicates that marital status has a noticeable influence on purchasing behavior.


```
In [34]: n_cols = 2
n_rows = -(~len(total_spent)) // n_cols

fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 4 * n_rows))
axes = axes.flatten()

for i, feature in enumerate(total_spent):
    sns.histplot(data=df, x=feature, hue='Children', kde=True, bins=30, ax=axes[i])
    axes[i].set_title(f"Distribution of {feature}")


plt.tight_layout()
plt.suptitle("Total Spent based on Children", fontsize=22, y=1.02)
plt.show()
```

## Total Spent based on Children




The histograms display spending across product categories according to the number of children in the household. Households with no children show a wider range of spending, including some very high spenders, while families with one or more children spend moderately, particularly on discretionary items like sweets and wine. This suggests that household size influences customer spending patterns.

```
In [35]: plt.figure(figsize=(6,4))
sns.boxplot(data=df, x='Children', y='Income', palette='Set1')
plt.title('Income by Number of Children')
plt.show()
```




```
In [36]: sns.scatterplot(x='Income', y='TotalSpent', data=df)  
plt.show()
```



This scatter plot displays the relationship between customers' income and their total spending. A positive trend can be observed: clients with higher income tend to spend more overall, although the relationship is not perfectly linear. Some high-income customers still show moderate spending, which may

indicate different spending habits or levels of engagement. The spread of points suggests significant variability in spending behavior across income levels, highlighting the need for segmentation to better understand these patterns.

```
In [37]: sns.boxplot(x='Children', y='TotalSpent', data=df)
plt.show()
```



The boxplot shows how total spending varies with the number of children in the household.

# Clustering Process

## 1. Scaling

```
In [38]: features = ['Age', 'Income', 'Children', 'Recency', 'TotalSpent']
```


The features were scaled using RobustScaler, this makes it more resistant to outliers observed in the income and spending distributions.

As a result, even though a few customers have unusually high incomes or spending levels, these outliers do not distort the clustering results, ensuring more reliable and stable cluster formation.

```
In [39]: from sklearn.preprocessing import RobustScaler  
scaler = RobustScaler()  
X_scaled = scaler.fit_transform(df[features])
```

## 2. Elbow Method

```
In [40]: from sklearn.cluster import KMeans  
import matplotlib.pyplot as plt  
  
list = []  
K = range(1, 11)  
  
for k in K:  
    kmeans = KMeans(n_clusters=k, random_state=42)  
    kmeans.fit(X_scaled)  
    list.append(kmeans.inertia_)  
  
plt.figure(figsize=(7,5))  
plt.plot(K, list, 'bo-')  
plt.xlabel('Number of clusters (k)')  
plt.ylabel('Inertia')  
plt.title('Elbow Method for Optimal k')  
plt.show()
```



Although the Elbow method did not show a clear inflection point, four clusters were retained as the most meaningful and actionable segmentation. This allows us to capture meaningful customer segments without creating too many small or overlapping clusters.

### 3. K-means

```
In [41]: best_k = 4
kmeans = KMeans(n_clusters=best_k, init='k-means++', random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)
```

## Cluster Analysis and Interpretation

```
In [42]: df['Cluster'].value_counts()
```

```
Out[42]: Cluster
0    677
1    600
2    488
3    447
Name: count, dtype: int64
```

```
In [43]: df.groupby('Cluster')[features].mean().round(2)
```

	Age	Income	Children	Recency	TotalSpent
Cluster					
<b>0</b>	47.63	31699.32	0.83	48.90	140.95
<b>1</b>	64.57	59476.98	0.89	47.01	741.61
<b>2</b>	54.74	79703.03	0.12	51.02	1447.72
<b>3</b>	59.49	43616.56	2.11	49.86	212.91

To describe each cluster, I calculated the mean of numerical features, which represents the typical value for a customer in that cluster.

Cluster 3: High-income, high-spending customers (Age: 54, Income: 79K, TotalSpent: 1,443).

These are the most valuable clients in terms of revenue per customer, with very few children. They are a top priority for loyalty programs and premium offers.

Cluster 1: Low-income, low-spending customers (Age: 47, Income: 32K, TotalSpent: 147).


Young and less engaged clients. They may respond well to targeted promotions to increase engagement.

Cluster 0: Older, medium-income customers (Age: 66, Income: 59K,

TotalSpent: 712).  
Stable, moderately spending clients. Important to retain these customers.

Cluster 2: Families with children (Age: 60, Income: 44K, Children: 2.1, TotalSpent: 220).  
Medium-income families, likely price-sensitive, with more discount purchases. They benefit from family-oriented offers.

```
In [44]: sns.scatterplot(data=df, x='Income', y='TotalSpent', hue='Cluster', palette='Set1')
plt.title('Clusters based on Income and Total Spending')
plt.show()
```



The scatterplot shows customers segmented into clusters based on their income and total spending.

```
In [45]: marital_distribution = pd.DataFrame()

for c in sorted(df['Cluster'].unique()):
    counts = df[df['Cluster']==c]['Marital_Status'].value_counts()
    marital_distribution[c] = counts

marital_distribution = marital_distribution.fillna(0).astype(int)

marital_distribution.columns = [f'Cluster {c}' for c in marital_distribution.col
marital_distribution
```

Out[45]:

Cluster 0 Cluster 1 Cluster 2 Cluster 3

Marital_Status		Cluster 0	Cluster 1	Cluster 2	Cluster 3
<b>Married</b>	278	226	176	177	
<b>Single</b>	173	96	122	83	
<b>Together</b>	168	157	128	120	
<b>Divorced</b>	50	85	46	51	
<b>Widow</b>	8	36	16	16	

The table shows that all marital statuses are distributed across the clusters without any cluster being dominated by a single category.

This indicates that marital status alone is not a strong discriminating feature for clustering customers, and other features such as income, spending, and purchasing behavior are more informative for segmenting the customer base.

In [46]: df.groupby('Cluster')[['NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchase

Out[46]: NumWebPurchases NumCatalogPurchases NumStorePurchases NumDealsPurch

Cluster	NumWebPurchases	NumCatalogPurchases	NumStorePurchases	NumDealsPurch
<b>0</b>	2.56	0.72	3.49	
<b>1</b>	5.76	3.39	7.64	
<b>2</b>	5.02	6.03	8.30	
<b>3</b>	3.13	0.98	4.11	

Analyzing purchase behavior across clusters revealed distinct shopping preferences.

- \* Cluster 0: customers buy through multiple channels with moderate web and catalog purchases and higher in-store purchases, showing balanced engagement and moderate responsiveness to promotions.
  - \* Cluster 1: members have low spending across all channels but visit the website frequently, indicating low conversion despite high online activity.
  - Cluster 2: consists of deal-sensitive families who make moderate purchases in-store and via catalog, showing high responsiveness to promotions.
  - \* Cluster 3: represents high-spending premium customers who buy across all channels but are less influenced by deals, purchasing efficiently with fewer visits.
- These insights complement the income and total spending segmentation, providing a detailed understanding of customer behavior for targeted marketing strategies.

### Business Insights:

Cluster 3 are top-value customers worth prioritizing for loyalty programs.  
Cluster 1 are low-engagement clients who may respond well to promotions.  
Cluster 2 could benefit from family-oriented offers or bundle discounts.  
Cluster 0 likely includes older, stable customers – focus on retention.

### Conclusion:

This project demonstrates how unsupervised learning (K-Means) can reveal meaningful customer segments. The identified clusters highlight differences in income, spending behavior, and engagement level, providing actionable insights for marketing and customer relationship management.